

Project work, part 3 - Data quality

Orie Kimura

November 7, 2025

General

links to your public GitHub repository and Streamlit app for the compulsory work.

<https://github.com/oriekimura123/IND320-Projectwork>

<https://oriekimura123-ind320.streamlit.app/>

Tasks

Accounts and repositories

Reuse your account, repository and Streamlit app from the previous part of the project work. Until peer review and feedback have been completed, push to a temporary GitHub branch for later merging.

API

Familiarise yourself with the API connection at <https://open-meteo.com/en/docs> til en ekstern side. Observe how you can select features and produce Python code. Be aware of multiple sub-sections that alter which type of data is selected.

Jupyter Notebook

- Use Oslo, Kristiansand, Trondheim, Tromsø and Bergen as representatives for the five electricity price areas in Norway. Find their geographical centre points in longitude and latitude. Save price area codes, city names, longitude and latitude in a Pandas data frame.

```
In [2]: import pandas as pd

city_data = {
    'pricearea':['N01', 'N02', 'N03', 'N04', 'N05'],
    'city_names':['Oslo', 'Kristiansand', 'Trondheim', 'Tromsø', 'Bergen'],
    'lon':[10.7461, 7.9956, 10.3951, 18.957, 5.328],
```

```

    'lat':[59.9127, 58.1467, 63.4305, 69.6496, 60.392]
}

df = pd.DataFrame(city_data)
df.head

```

```

Out[2]: <bound method NDFrame.head of      pricearea  city_names  lon  lat
0      N01      Oslo  10.7461  59.9127
1      N02  Kristiansand  7.9956  58.1467
2      N03    Trondheim  10.3951  63.4305
3      N04    Tromsø  18.9570  69.6496
4      N05    Bergen  5.3280  60.3920>

```

- Use the open-meteo API to retrieve historical reanalysis data using the ERA5 model for a single location as follows:
 - Create a function for the API download task that takes a pair of longitude and latitude values, plus a year as input and downloads the same weather properties as were used in the CSV file in part 1 of the project work.

```

In [5]: # Function for the API download task that takes a pair of Longitude and Latitude v
# and downlods the same weather properties
# temperature_2m (°C)    precipitation (mm)          wind_speed_10m (m/s)    wind_gusts_

import requests
import pandas as pd

def get_era5_data(latitude: float, longitude: float, year: str) -> pd.DataFrame | N
    """
    Get historical ERA5 weather data from Open-Meteo.

    Args:
        latitude (float): Latitude coordinate (e.g., 52.52).
        longitude (float): Longitude coordinate (e.g., 13.41).
        year (str): Start date in 'YYYY' format (e.g., '2019').

    Returns:
        pd.DataFrame or None: A Pandas DataFrame with the hourly weather data,
                               or None if the API request fails.

    """

    # Define the API endpoint and the required hourly variables
    BASE_URL = "https://archive-api.open-meteo.com/v1/era5"

    start_date = year + "-01-01"
    end_date = year + "-12-31"

    # The properties
    Properties = [
        "temperature_2m",
        "precipitation",
        "wind_speed_10m",
        "wind_gusts_10m",
        "wind_direction_10m"
    ]

```

```

# Construct the parameters dictionary
params = {
    "latitude": latitude,
    "longitude": longitude,
    "start_date": start_date,
    "end_date": end_date,
    "hourly": Properties,
    "timezone": "auto",
    "models": "era5",
    "wind_speed_unit": "ms"
}

# Make the API request
try:
    response = requests.get(BASE_URL, params=params)
    response.raise_for_status() # Raise an exception for bad status codes (4xx)
    data = response.json()
except requests.exceptions.RequestException as e:
    print(f"Error fetching data: {e}")
    return None

# Process the JSON response into a Pandas DataFrame
# The 'hourly' key contains the time series data
hourly_data = data.get("hourly", {})

if not hourly_data:
    print("API response is missing 'hourly' data.")
    return None

df = pd.DataFrame(hourly_data)

# Convert the 'time' column to a datetime index
df['time'] = pd.to_datetime(df['time'])
df = df.set_index('time')

return df

```

- Apply the function to download data for Bergen for the year 2019.

```

In [6]: # Bergen
LAT = 60.392
LON = 5.328

YEAR = "2019"

# Call the function
df_Bergen = get_era5_data(LAT, LON, YEAR)

if df_Bergen is not None:
    print("Bergens hourly weather data, 2019")
    print(df_Bergen.head())
    print(f"\nDataFrame Shape: {df_Bergen.shape}")

```

Bergens hourly weather data, 2019

time	temperature_2m	precipitation	wind_speed_10m \
2019-01-01 00:00:00	7.2	0.3	10.95
2019-01-01 01:00:00	7.1	0.4	12.20
2019-01-01 02:00:00	7.3	0.5	13.08
2019-01-01 03:00:00	7.5	0.7	14.01
2019-01-01 04:00:00	7.0	0.5	14.29

time	wind_gusts_10m	wind_direction_10m
2019-01-01 00:00:00	24.2	265
2019-01-01 01:00:00	26.4	280
2019-01-01 02:00:00	23.7	293
2019-01-01 03:00:00	25.5	309
2019-01-01 04:00:00	28.5	314

DataFrame Shape: (8760, 5)

Outliers and anomalies:

- Plot the temperature as a function of time.
 - Perform a high-pass filtering of the temperature using Direct Cosine Transfer to create seasonally adjusted temperature variations (SATV).
 - Add curves to the plot indicating Statistical Process Control boundaries between inliers and outliers based on the SATV according to robust statistics estimated from the whole year. Colour outliers with a contrasting colour. Do not plot SATV values; only use them to find boundaries and outliers.
 - Let the frequency cut-off for the DCT and the number of standard deviations be parameters with sensible defaults.
 - Wrap this in a function that returns the plot and relevant summaries of the outliers, and test the function.

```
In [7]: # Function to
# Perform a high-pass filtering of the temperature using Direct Cosine Transfer
# to create seasonally adjusted temperature variations (SATV).

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.fft import dct, idct
from typing import Tuple, Dict, Any

def calculate_satv(temperature_series: pd.Series, dct_cutoff: int = 3) -> pd.Series:
    """
    Performs high-pass filtering using DCT to calculate Seasonally Adjusted
    Temperature Variations (SATV), removing the first 'cutoff_count' DCT coefficients.
    """
    T = temperature_series.values
    N = len(T)

    # Apply Forward DCT
```

```

T_dct = dct(T, type=1, norm='forward')

# Define High-Pass Filter Coefficients
T_dct_hp = T_dct.copy()

if dct_cutoff >= 0 and dct_cutoff <= N:
    # Zero out the Low-frequency coefficients (k=0 up to cutoff_count-1)
    T_dct_hp[:dct_cutoff] = 0
else:
    raise ValueError("cutoff_count must be between 0 and the length of the time

# Apply Inverse DCT
SATV_values = idct(T_dct_hp, type=1, norm='forward')

# Return as a Pandas Series
return pd.Series(SATV_values, index=temperature_series.index)

```

```

In [8]: # Add curves to the plot indicating Statistical Process Control boundaries between
# based on the SATV according to robust statistics estimated from the whole year.
# Colour outliers with a contrasting colour.
# The frequency cut-off for the DCT and the number of standard deviations be paramete

def plot_anomaly_control(
    df: pd.DataFrame,
    temp_col: str = 'temperature_2m',
    dct_cutoff: int = 3,
    std_dev_multiple: float = 3.0,
    figsize: Tuple[int, int] = (12, 6)
) -> Tuple[plt.Figure, Dict[str, Any]]:

    """
    Plots the original temperature data with Statistical Process Control boundaries
    based on Seasonally Adjusted Temperature Variations (SATV).
    """

    if temp_col not in df.columns:
        raise ValueError(f"Column '{temp_col}' not found in DataFrame.")

    # Calculate SATV
    temp_series = df[temp_col]
    # Pass the renamed parameter
    satv_series = calculate_satv(temp_series, dct_cutoff=dct_cutoff)

    # Estimate Robust Statistics from SATV
    median = satv_series.median()
    mad = (satv_series - median).abs().median()
    robust_std = mad * 1.4826

    # Define the control/SATV Limits
    upper_satv_limit = median + std_dev_multiple * robust_std
    lower_satv_limit = median - std_dev_multiple * robust_std

    # Identify Outliers
    outlier_mask = (satv_series > upper_satv_limit) | (satv_series < lower_satv_lim
    outliers = df[outlier_mask]

    # Calculate Control Boundaries for the Original Temperature Data

```

```

seasonal_cycle = temp_series - satv_series
upper_boundary = seasonal_cycle + upper_satv_limit
lower_boundary = seasonal_cycle + lower_satv_limit

# Plott the results
fig, ax = plt.subplots(figsize=figsize)

# Plot the original temperature data (inliers)
ax.plot(
    df.index, df[temp_col],
    label=f'Original {temp_col}',
    color='blue',
    linewidth=1
)

# Plot the control boundaries
ax.plot(upper_boundary, label=f'{std_dev_multiple} MAD Upper Limit', color='red')
ax.plot(lower_boundary, label=f'{std_dev_multiple} MAD Lower Limit', color='red')

# Colour outliers
ax.scatter(
    outliers.index, outliers[temp_col],
    color='gold',
    edgecolor='black',
    s=40,
    zorder=5,
    label='Outliers'
)

# Set plot aesthetics (Updated title to reflect the new parameter name)
ax.set_title(
    f'Temperature Outlier Detection (DCT Cutoff  $\rightarrow$  {dct_cutoff}, \
    {std_dev_multiple}  $\times$  MAD)', fontsize=12)
ax.set_xlabel('Time')
ax.set_ylabel(f'{temp_col} (°C)')
ax.grid(True, linestyle='--', alpha=0.7)
ax.legend()

# Prepare Summary
summary = {
    'total_observations': len(df),
    'number_of_outliers': len(outliers),
    'percentage_outliers': (len(outliers) / len(df)) * 100,
    'robust_median': median,
    'robust_std_estimate': robust_std,
    'outlier_data': outliers[[temp_col]]
}

return fig, summary

```

In [9]: # --- Setup Example Data for Testing (Requires df_Bergen to be Loaded) ---

```

dates = pd.date_range('2019-01-01', periods=365*24, freq='h')
t_index = np.arange(365*24)
seasonal_cycle = 10 * np.sin(2 * np.pi * t_index / (365*24)) # Annual cycle
noise = np.random.normal(0, 1.5, len(dates)) # Normal variation

```

```

# Add an extreme outlier for testing
noise[1000:1010] += 10
noise[2000:2050] -= 5
noise[4500:4510] -= 7
noise[6540:6550] += 5

temperature = 5 + seasonal_cycle + noise
df_Bergen_tmp = pd.DataFrame({'temperature_2m': temperature}, index=dates)

```

```

In [10]: # --- Test 1
selected_dct_cutoff=3
selected_std_dev_multiple = 3.0

print("--- Test 1:")
plot1, summary1 = plot_anomaly_control(
    df_Bergen_tmp,
    temp_col = 'temperature_2m',
    dct_cutoff = selected_dct_cutoff,
    std_dev_multiple = selected_std_dev_multiple
)

print("\nSummary 1:")
print(f"Total Observations: {summary1['total_observations']}")
print(f"Number of Outliers: {summary1['number_of_outliers']}")
print(f"Percentage Outliers: {summary1['percentage_outliers']:.2f}%")
print(f"Robust Median of SATV: {summary1['robust_median']:.2f}")
print("First 5 Outlier Records:")
print(summary1['outlier_data'].head())

plt.show() # Uncomment to display the plot

```

```

findfont: Font family ['cmsy10'] not found. Falling back to DejaVu Sans.
findfont: Font family ['cmr10'] not found. Falling back to DejaVu Sans.
findfont: Font family ['cmtt10'] not found. Falling back to DejaVu Sans.
findfont: Font family ['cmi10'] not found. Falling back to DejaVu Sans.
findfont: Font family ['cmb10'] not found. Falling back to DejaVu Sans.
findfont: Font family ['cmss10'] not found. Falling back to DejaVu Sans.
findfont: Font family ['cmex10'] not found. Falling back to DejaVu Sans.


```

--- Test 1:

```

Summary 1:
Total Observations: 8760
Number of Outliers: 0
Percentage Outliers: 0.00%
Robust Median of SATV: -0.07
First 5 Outlier Records:
Empty DataFrame
Columns: [temperature_2m]
Index: []

```



```
In [11]: # --- Test 2
selected_dct_cutoff = 3
selected_std_dev_multiple = 2.0

print("\n--- Test 2")
plot2, summary2 = plot_anomaly_control(
    df_Bergen_tmp,
    temp_col = 'temperature_2m',
    dct_cutoff = selected_dct_cutoff,
    std_dev_multiple = selected_std_dev_multiple)

print("\nSummary 2:")
print(f"Total Observations: {summary2['total_observations']}")
print(f"Number of Outliers: {summary2['number_of_outliers']}")
print(f"Percentage Outliers: {summary2['percentage_outliers']:.2f}%")
print(f"Robust Median of SATV: {summary2['robust_median']:.2f}")
print("First 5 Outlier Records:")
print(summary2['outlier_data'].head())
plt.show() # Uncomment to display the second plot
```

--- Test 2

Summary 2:

Total Observations: 8760


Number of Outliers: 45

Percentage Outliers: 0.51%

Robust Median of SATV: -0.07

First 5 Outlier Records:

	temperature_2m
2019-01-01 00:00:00	1.331580
2019-01-01 04:00:00	3.260465
2019-01-01 08:00:00	3.340842
2019-01-01 12:00:00	3.712739
2019-01-01 16:00:00	3.770238



- Plot the precipitation as a function of time.
 - Indicate anomalies according to the Local Outlier Factor method.
 - Let the proportion of outliers be a parameter defaulting to 1%.
 - Wrap this in a function that returns the plot and relevant summaries of the outliers, and test the function.

```
In [12]: import matplotlib.pyplot as plt

print(df_Bergen.head())


plt.figure(figsize=(10,5))

df_Bergen['precipitation'].plot(
    title='Precipitation, Bergen, 2019',
    ylabel='precipitation (mm)'
)

plt.grid()
plt.show()
```

	temperature_2m	precipitation	wind_speed_10m \
time			
2019-01-01 00:00:00	7.2	0.3	10.95
2019-01-01 01:00:00	7.1	0.4	12.20
2019-01-01 02:00:00	7.3	0.5	13.08
2019-01-01 03:00:00	7.5	0.7	14.01
2019-01-01 04:00:00	7.0	0.5	14.29

	wind_gusts_10m	wind_direction_10m
time		
2019-01-01 00:00:00	24.2	265
2019-01-01 01:00:00	26.4	280
2019-01-01 02:00:00	23.7	293
2019-01-01 03:00:00	25.5	309
2019-01-01 04:00:00	28.5	314



```
In [13]: # Function
# to calculate anomalies according to the Local Outlier Factor method
# to summarize the outliers
# plot the result

import numpy as np
from sklearn.neighbors import LocalOutlierFactor
import matplotlib.pyplot as plt
from typing import Tuple, Dict, Any

def calculate_LOF_anomalies(
    df_full: pd.DataFrame,
    feature_col: str,
    n_neighbors: int,
    contamination: float
) -> Tuple[plt.Figure, pd.DataFrame]:

    """
```

Calculate anomalies according to the Local Outlier Factor method,
summarize of the outliers and plot the result

Args:

`df_full` (pd.DataFrame): The full DataFrame containing the feature to analyze and all context columns (e.g., 'temperature_2m').
`feature_col` (str): The name of the column to run LOF on (e.g., 'precipitation').
`n_neighbors` (int): Number of neighbors to use for LOF.
`contamination` (float): The proportion of outliers expected in the dataset.

Returns:

Tuple[plt.Figure, pd.DataFrame]: The generated matplotlib Figure and a DataFrame summarizing the outliers.

"""

Filter out potential NaNs for robust LOF calculation

`df_clean = df_full.dropna(subset=[feature_col]).copy()`

Run LOF Analysis

`lof = LocalOutlierFactor(n_neighbors=n_neighbors, contamination=contamination)`
`pred_labels = lof.fit_predict(X=df_clean)`

Store scores and labels back into the clean DataFrame

`df_clean['LOF_Label'] = pred_labels`
`df_clean['LOF_Score'] = np.clip(lof.negative_outlier_factor_, -10, 0)`

Separate Inliers/Outliers and Calculate Baseline

`inliers = df_clean[df_clean['LOF_Label'] == 1]`
`outliers = df_clean[df_clean['LOF_Label'] == -1]`

`median_baseline = df_clean[feature_col].median()`
`mean_baseline = df_clean[feature_col].mean()`

for Statistical Summary

`total_samples = len(df_clean)`
`num_outliers = len(outliers)`
`outlier_percentage = (num_outliers / total_samples) * 100`

`stats_summary = {`
 `'total_samples': total_samples,`
 `'num_outliers': num_outliers,`
 `'outlier_percentage': f'{outlier_percentage:.2f}%' # Format to 2 decimal places`
`}`

Generate Plot

`fig, ax = plt.subplots(figsize=(10, 6))`

Plot the inliers using the time index

`ax.plot(inliers.index, inliers[feature_col], 'o', alpha=0.5, label='Inlier')`

Plot the outliers using the time index

`ax.plot(outliers.index, outliers[feature_col], 'o', color='red', label='Outlier')`

Plot Aesthetics

`ax.set_xlabel('Time')`

```

ax.set_ylabel(f'{feature_col}')
ax.axhline(mean_baseline, color='gray', linestyle=':', label='Mean Baseline')
ax.legend(loc='upper right')
ax.set_title(f'{feature_col.capitalize()} Anomaly Detection using LOF (n_neighb
ax.grid(True, linestyle='--', alpha=0.6)

plt.tight_layout() # Adjust layout to prevent clipping

# Prepare Summary DataFrame
# Use localized formatting to display clean output
pd.set_option('display.float_format', '{:.2f}'.format)

outlier_summary = outliers.copy()

# Add Baseline Metrics
outlier_summary['Median_Baseline'] = median_baseline
outlier_summary['Absolute_Deviation'] = (outlier_summary[feature_col] - mean_ba

# Select and rename final columns for the summary
summary_cols = {
    feature_col: feature_col.capitalize(),
    'LOF_Score': 'LOF_Score',
    'Median_Baseline': 'Median_Baseline',
    'Absolute_Deviation': 'Absolute_Deviation',
}

# Add temperature and wind if they exist in the original DataFrame
if 'temperature_2m' in df_full.columns:
    summary_cols['temperature_2m'] = 'Temperature'
if 'wind_gusts_10m' in df_full.columns:
    summary_cols['wind_gusts_10m'] = 'Wind_Gusts'

# Final selection and sorting (by LOF_Score for most severe first)
outlier_summary_final = outlier_summary[summary_cols.keys()] \
    .rename(columns=summary_cols) \
    .sort_values(by='LOF_Score')

return fig, outlier_summary_final, stats_summary

```

In [14]: # Test 1 (Requires df_Bergen to be loaded) ---

```

test1_neighbors=20
test1_contamination=0.01
fig1, summary1, stats1 = calculate_LOF_anomalies(
    df_Bergen,
    feature_col="precipitation",
    n_neighbors=test1_neighbors,
    contamination=test1_contamination)

print(f"\n--- TEST 1 (n_neighbors={test1_neighbors}, Contamination={test1_contamina

print("Statistics:", stats1)
print("Outlier Summary Head:\n", summary1.head())

```


```
--- TEST 1 (n_neighbors=20, Contamination=0.01) ---
```

```
Statistics: {'total_samples': 8760, 'num_outliers': 88, 'outlier_percentage': '1.00%'}
```

```
Outlier Summary Head:
```

	Precipitation	LOF_Score	Median_Baseline \
time			
2019-01-01 04:00:00	0.50	-2.60	0.00
2019-01-01 06:00:00	0.40	-2.59	0.00
2019-01-01 05:00:00	1.30	-2.53	0.00
2019-01-01 03:00:00	0.70	-2.38	0.00
2019-01-01 07:00:00	0.50	-2.35	0.00

	Absolute_Deviation	Temperature	Wind_Gusts
time			
2019-01-01 04:00:00	0.26	7.00	28.50
2019-01-01 06:00:00	0.16	6.50	34.70
2019-01-01 05:00:00	1.06	6.50	34.30
2019-01-01 03:00:00	0.46	7.50	25.50
2019-01-01 07:00:00	0.26	5.90	31.20



```
In [15]: # Test 2 (Requires df_Bergen to be Loaded) ---
```

```
test2_neighbors=20
test2_contamination=0.05
fig2, summary2, stats2 = calculate_LOF_anomalies(
    df_Bergen,
    feature_col="precipitation",
    n_neighbors=test2_neighbors,
    contamination=test2_contamination)

print(f"\n--- TEST 2 (n_neighbors={test2_neighbors}, Contamination={test2_contamina

print("Statistics:", stats2)
print("Outlier Summary Head:\n", summary2.head())
```


--- TEST 2 (n_neighbors=20, Contamination=0.05) ---

Statistics: {'total_samples': 8760, 'num_outliers': 438, 'outlier_percentage': '5.00%'}

Outlier Summary Head:

	Precipitation	LOF_Score	Median_Baseline \
time			
2019-01-01 04:00:00	0.50	-2.60	0.00
2019-01-01 06:00:00	0.40	-2.59	0.00
2019-01-01 05:00:00	1.30	-2.53	0.00
2019-01-01 03:00:00	0.70	-2.38	0.00
2019-01-01 07:00:00	0.50	-2.35	0.00

	Absolute_Deviation	Temperature	Wind_Gusts
time			
2019-01-01 04:00:00	0.26	7.00	28.50
2019-01-01 06:00:00	0.16	6.50	34.70
2019-01-01 05:00:00	1.06	6.50	34.30
2019-01-01 03:00:00	0.46	7.50	25.50
2019-01-01 07:00:00	0.26	5.90	31.20



Seasonal-Trend decomposition using LOESS (STL)

- Perform LOESS on the production data from elbub (downloaded in part 1 of the project) and plot its decomposition.
 - Let the electricity price area, production group, period length, seasonal smoother, trend smoother and robust (true/false) be parameters, and give each of them sensible defaults.
 - Wrap this in a function that returns the plot, and test the function.

```
In [16]: import pandas as pd
import json
import os
```

```

def load_elbub_data(filepath='elbub_data.json'):
    """
    Loads nested elbub JSON data, which is a list of dictionaries,
    extracting production data from ALL available price areas.
    """
    try:
        with open(filepath, 'r') as f:
            data = json.load(f)

        all_data_frames = []

        # 1. Iterate through the top-level list
        # Each item in 'data' is a dictionary corresponding to a block (like NO1, N
        for block in data:
            # Check if the block is the one containing time-series attributes
            if 'attributes' in block and 'productionPerGroupMbaHour' in block['attr

                attributes = block['attributes']

                # Extract the price area name for filtering later
                area_name = attributes.get('name')

                # Extract the list of time-series records
                production_list = attributes['productionPerGroupMbaHour']

                # Create a temporary DataFrame for this block
                df_area = pd.DataFrame(production_list)

                # Add the area name to all rows for filtering/grouping
                df_area['priceArea'] = area_name

                all_data_frames.append(df_area)

        if not all_data_frames:
            raise ValueError("No production data blocks found in the JSON file.")

        # 2. Concatenate all price areas into a single DataFrame
        df_full = pd.concat(all_data_frames, ignore_index=True)

        # 3. Clean and process the combined DataFrame
        df_full['startTime'] = pd.to_datetime(df_full['startTime'], utc=True)
        df_full = df_full.set_index('startTime').sort_index()

        # Ensure quantity is numeric
        df_full['quantityKwh'] = pd.to_numeric(df_full['quantityKwh'], errors='coer

        return df_full

    except FileNotFoundError:
        print(f"Error: The file '{filepath}' was not found.")
        return pd.DataFrame()
    except Exception as e:
        print(f"Error loading elbub data. Check file structure. Error: {e}")
        return pd.DataFrame()

```

```
repo_root = os.path.abspath(os.path.join(os.path.dirname(''), '..'))
print (repo_root)
DOWNLOADS_FOLDER = os.path.join(repo_root, 'downloads')
filename = 'Production_per_group_mad_hours_2021.json'
FILE_PATH = DOWNLOADS_FOLDER + '\\' + filename

df_elbub = load_elbub_data(FILE_PATH)
if not df_elbub.empty:
    print(df_elbub.head())
    print(df_elbub.tail())
```

C:\Users\oriek\IND320\IND320-Projectwork-Orie3

endTime \

startTime

2020-12-31 23:00:00+00:00	2021-01-01T01:00:00+01:00
2020-12-31 23:00:00+00:00	2021-01-01T01:00:00+01:00
2020-12-31 23:00:00+00:00	2021-01-01T01:00:00+01:00
2020-12-31 23:00:00+00:00	2021-01-01T01:00:00+01:00
2020-12-31 23:00:00+00:00	2021-01-01T01:00:00+01:00

lastUpdatedTime priceArea \

startTime

2020-12-31 23:00:00+00:00	2024-12-20T10:35:40+01:00	N01
2020-12-31 23:00:00+00:00	2024-12-20T10:35:40+01:00	N03
2020-12-31 23:00:00+00:00	2024-12-20T10:35:40+01:00	N01
2020-12-31 23:00:00+00:00	2024-12-20T10:35:40+01:00	N05
2020-12-31 23:00:00+00:00	2024-12-20T10:35:40+01:00	N03

productionGroup quantityKwh

startTime

2020-12-31 23:00:00+00:00	hydro	2507716.80
2020-12-31 23:00:00+00:00	solar	19.72
2020-12-31 23:00:00+00:00	solar	6.11
2020-12-31 23:00:00+00:00	thermal	77742.00
2020-12-31 23:00:00+00:00	hydro	2836774.00

endTime \

startTime

2021-12-31 22:00:00+00:00	2022-01-01T00:00:00+01:00
2021-12-31 22:00:00+00:00	2022-01-01T00:00:00+01:00
2021-12-31 22:00:00+00:00	2022-01-01T00:00:00+01:00
2021-12-31 22:00:00+00:00	2022-01-01T00:00:00+01:00
2021-12-31 22:00:00+00:00	2022-01-01T00:00:00+01:00

lastUpdatedTime priceArea \

startTime

2021-12-31 22:00:00+00:00	2024-10-27T07:56:27+01:00	N01
2021-12-31 22:00:00+00:00	2024-10-27T07:56:27+01:00	N04
2021-12-31 22:00:00+00:00	2024-10-27T07:56:27+01:00	N01
2021-12-31 22:00:00+00:00	2024-10-27T07:56:27+01:00	N03
2021-12-31 22:00:00+00:00	2024-10-27T07:56:27+01:00	N05

productionGroup quantityKwh

startTime

2021-12-31 22:00:00+00:00	other	0.40
2021-12-31 22:00:00+00:00	other	46.12
2021-12-31 22:00:00+00:00	hydro	1286038.20
2021-12-31 22:00:00+00:00	other	80.59
2021-12-31 22:00:00+00:00	wind	0.00

```
In [17]: import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import STL
from typing import Tuple, Dict, Any

def decompose_elbub_production(
    df: pd.DataFrame,
    area: str = 'N03',
```

```

prod_group: str = 'hydro',
period: int = 24,          # Sensible default: 24 hours (daily seasonality)
seasonal_smoother: int = 13, # Sensible default: odd number > 7, controls seasonality
trend_smoother: int = 501, # Sensible default: Large odd number for annual trend
robust: bool = True,      # Sensible default: True (robust to outliers/anomalies)
prod_col: str = 'quantityKwh'
) -> plt.Figure:
"""
    Performs STL decomposition (LOESS) on electricity production data and plots the decomposition.

    Args:
        df (pd.DataFrame): The full DataFrame loaded from elbub JSON.
        area (str): Electricity price area (e.g., 'N01').
        prod_group (str): Production group (e.g., 'hydro').
        period (int): Length of the seasonal cycle (e.g., 24 for daily, 24*7 for weekly).
        seasonal_smoother (int): Odd span (window size) for seasonal smoothing (must be odd).
        trend_smoother (int): Odd span for trend smoothing (must be odd).
        robust (bool): If True, uses robust fitting (less sensitive to outliers).
        prod_col (str): The name of the production column.

    Returns:
        plt.Figure: The Matplotlib Figure object containing the decomposition plot.
"""

# Price Area filter
area_filter = df['priceArea'].str.strip() == area

# prod_group filter
prod_group_filter = df['productionGroup'].str.strip() == prod_group

# Apply the combined filter
df_filtered = df[area_filter & prod_group_filter].copy()

# Check for emptiness
if df_filtered.empty:
    # Raise an informative error based on the original input
    raise ValueError(f"No data found for Area: {area} and Group: {prod_group}").

# Ensure the Series is continuous and sorted
series = df_filtered[prod_col].sort_index()

# Check if the series is long enough
min_length = period + 1
if len(series) < min_length:
    raise ValueError(f"Series too short for period {period}. Requires at least {min_length} data points.")

try:
    # The key parameters match the function arguments:
    res = STL(
        series,
        period=period,
        seasonal=seasonal_smoother,
        trend=trend_smoother,
        robust=robust
    ).fit()
except Exception as e:

```

```

        print(f"STL decomposition failed. Check period and smoother spans: {e}")
        return None

    # Plot the Decomposition
    fig = plt.figure(figsize=(20, 30))
    res.plot(fig)

    for ax in fig.axes:
        # Rotate x-axis labels by 45 degrees for better readability
        ax.tick_params(axis='x', rotation=45)

    # Customizing the Plot Title for clarity
    fig.suptitle(
        f'STL Decomposition of {prod_group} Production in {area}, 2021',
        fontsize=12,
        y=1.02
    )

    # Add parameter summary to the plot (optional, but helpful for testing)
    fig.text(0.02, 0.03,
             f'P={period}, S_span={seasonal_smoother}, T_span={trend_smoother}, Rob
             fontsize=8, alpha=0.7, transform=fig.transFigure)

    plt.tight_layout()
    plt.tight_layout(rect=[0, 0.05, 1, 0.98])

    return fig

```

```

In [18]: # Test 1: Using all default parameters (Daily seasonality)
daily = 24
period = daily
selected_area = 'N01'
selected_prod_group = 'hydro'


print("--- Test 1 ---")
plot1 = decompose_elbub_production(
    df_elbub,
    area = selected_area,
    prod_group = selected_prod_group,
    period=period,
    robust=False)
plt.show() # Uncomment to display the plot

```

--- Test 1 ---

<Figure size 2000x3000 with 0 Axes>

STL Decomposition of hydro Production in NO1, 2021



P=24, S_span=13, T_span=501, Robust=False

```
In [19]: # Test 2
daily = 24
period = daily
selected_area = 'N02'
selected_prod_group = 'hydro'


print(f"\n--- TEST 2 ---")

plot2 = decompose_elbub_production(
    df_elbub,
    area = selected_area,
    prod_group = selected_prod_group,
    period=period,
    robust=True)
plt.show()
```

--- TEST 2 ---

<Figure size 2000x3000 with 0 Axes>

STL Decomposition of hydro Production in NO2, 2021



P=24, S_span=13, T_span=501, Robust=True

```
In [20]: # Test 3: Adjusting parameters
weekly = 24*7
period = weekly
selected_area = 'NO2'
selected_prod_group = 'hydro'


plot3 = decompose_elbub_production(
    df_elbub,
    period=period,
    area = selected_area,
    prod_group = selected_prod_group,
    robust=True
)

print(f"\n--- TEST 3 ---")
plt.show()
```

--- TEST 3 ---

<Figure size 2000x3000 with 0 Axes>

STL Decomposition of hydro Production in NO2, 2021



P=168, S_span=13, T_span=501, Robust=True

Spectrogram

- Create a spectrogram based on the production data from elhub.
 - Let the electricity price area, production group, window length and window overlap be parameters, and give each of them sensible defaults.
 - Wrap this in a function that returns the plot, and test the function.

```
In [28]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from scipy.signal import stft
from typing import Tuple

def plot_spectrogram_decomposition(
    df_elhub: pd.DataFrame,
    area: str,
    prod_group: str,
    window_length: int = 168,
    window_overlap: int = 84,
    figsize: Tuple[float, float] = (10, 6)
) -> plt.Figure:
    """
```

Filters the Elhub data and generates a two-panel plot:

1. Time Series of the filtered data.
2. Spectrogram (STFT) of the filtered data.

Args:

df_elhub (pd.DataFrame): The combined Elhub data DataFrame.
area (str): The price area (e.g., 'N03').
prod_group (str): The production group (e.g., 'hydro').
window_length (int): The STFT window length (nperseg). Defaults to 168 (weeks).
figsize (Tuple[float, float]): Size of the output figure.

Returns:

plt.figure: The Matplotlib Figure object.

"""

Setup and Filtering

df_spectrogram = df_elhub.copy()

window_overlap = window_length // 2 # noverlap is 50% of nperseg

Filtering with whitespace-removed

area_filter = df_spectrogram['priceArea'].str.strip() == area

prod_group_filter = df_spectrogram['productionGroup'].str.strip() == prod_group

Apply the combined filter

df_filtered2 = df_spectrogram[area_filter & prod_group_filter].copy()

Basic check to prevent crash if data is empty

if df_filtered2.empty:

print(f"Error: No data found for Area: {area} and Group: {prod_group}. Cannot")

return plt.figure()

--- 1. Perform STFT ---

f, t, Zxx = stft(
 df_filtered2['quantityKwh'],
 fs = 1.0, # Data is hourly (1 sample per hour)
 nperseg = window_length,
 noverlap = window_overlap
)

--- 2. Map STFT times back to actual datetimes (using the time vector t) ---

The time vector 't' is measured in hours (samples) since fs=1.0

start_time = df_filtered2.index[0]

The interval between samples is 1 hour (1 day / 24 samples)

time_step = pd.Timedelta(hours=1)

Map the numerical time vector 't' (in hours) back to datetime objects

*# datetime_labels = [start_time + i * time_step for i in t]*

start_time = df_filtered2.index[0]

time_step = pd.Timedelta(hours=1)

IMPORTANT FIX: Use the short time vector (t_stft) from the STFT result

datetime_labels = [start_time + pd.Timedelta(hours=i) **for** i **in** t]

--- 3. Plotting ---

fig, axs = plt.subplots(2, 1, figsize=figsize, sharex=**True**)

```

# Top Plot (Time Series)
axs[0].plot(df_filtered2.index, df_filtered2['quantityKwh'], label='Production')
axs[0].set_ylabel('Amplitude (QuantityKwh)')
axs[0].set_title(f'Time Series and Spectrogram for {prod_group} in {area} (Wind
axs[0].grid(True, linestyle='--', alpha=0.6)

# Bottom Plot (Spectrogram)
pcm = axs[1].pcolormesh(
    datetime_labels,
    f,
    np.abs(Zxx),
    shading='auto',
    vmin=0,
    vmax=np.max(np.abs(Zxx)) / 2
)
axs[1].set_ylabel('Frequency [cycles/hour]')
axs[1].set_xlabel('Time')
fig.colorbar(pcm, ax=axs[1], label='Amplitude')

# Rotate X-Labels for better readability
for ax in axs:
    ax.tick_params(axis='x', rotation=45)

plt.tight_layout()
return fig

```


```

In [29]: # Test 1
selected_area = 'N01'
selected_prod_group = 'hydro'
selected_window_length = 24
selected_window_overlap = 20

print("--- Test 1: DailySpectrogram ---")
fig_daily = plot_spectrogram_decomposition(
    df_elbub,
    area = selected_area,
    prod_group = selected_prod_group,
    window_length = selected_window_length,
    window_overlap = selected_window_overlap)
plt.show()

```


--- Test 1: DailySpectrogram ---



```
In [30]: # Test 2
selected_area = 'NO1'
selected_prod_group = 'hydro'
selected_window_length = 24 * 7
selected_window_overlap = selected_window_length // 2

print("--- Test 2: Weekly spectrogram---")
fig_weekly = plot_spectrogram_decomposition(
    df_elbub,
    area = selected_area,
    prod_group = selected_prod_group,
    window_length = selected_window_length,
    window_overlap = selected_window_overlap)
plt.show()
```

--- Test 2: Weekly spectrogram---



Tasks, Streamlit app

- Update your Streamlit app from part 2 of the project according to the following points.
- Move page 4 (the one with the price area selector) in front of page 2, and add a new page between page 4 and page 3 and a new page between page 3 and page 5, i.e., Old order: 1, 2, 3, 4, 5, new order: 1, 4, new A, 2, 3, new B, 5.
- Exchange the CSV import of meteorological data with the open-meteo API import.
 - Let the choice of downloaded data depend on the selector that is now on page 2. Let the chosen year be 2021.
 - Note! For each of the new pages below, consider whether you can depend on the area selector on page 2 or need a local one.
- On page "new A", use st.tabs() and fill the first tab with the STL analysis and the second tab with the Spectrogram. Add necessary UI elements and plots to both.
- On page "new B", use st.tabs() and fill the first tab with the Outlier/SPC analysis and the second tab with the Anomaly/LOF analysis. Add necessary UI elements, plots and statistics to both.

My Streamlit app is <https://oriekimura123-ind320.streamlit.app/>

Log Describing the Compulsory Work

Coding in Jupyter Notebook

I copied the required tasks into the notebook in Markdown format. Then I wrote and tested the code step by step, using help from AI tools when needed.

Once all the tasks were initially completed, I reviewed every cell to design how data is passed between functions so that everything works smoothly in the Streamlit app using `st.session_state`.

Finally, I organized the functions into separate Python files inside the `utils` folder.

Summary of Analysis Methods

Once the coding was done, I realized it was important to understand the purpose of each statistical method (what it identifies and analyzes) and how to interpret the results.

I used Gemini and ChatGPT to research each method and summarize their application within the context of my project (weather data and Elhub production data)

Coding the Streamlit App

One main challenge in the Streamlit app was ensuring smooth data transfer between different functions and pages. After some debugging, it worked relatively well. I also added more information about the app and each page on the front page for easier navigation.

Commit and Push to GitHub / Streamlit Cloud

I had to adjust the code to correctly handle absolute and relative paths, ensuring consistent functionality between my local environment and the cloud deployment

Brief Description of AI Usage

AI tools were used to generate code, fix bugs, and summarize technical concepts.

They also helped me learn and understand the statistical analysis methods applied in my project.

Generate initial versions of the code, improve and customize the code

I described what I needed, and the AI tools suggested initial drafts. Sometimes the code worked immediately, and other times I had to adjust it. After a few iterations, everything worked as expected.

Debug errors

I frequently consulted the AI tools to adjust my Streamlit app code for cloud deployment, especially concerning the use of functions from the `utils` folder. The AI provided various recommendations, including detailed `sys.path` manipulation to resolve import errors.

The path lines I chose for the main program (`My_Homepage.py`) look messy, but they are currently the only working solution for importing functions and ensuring deployment

consistency.

Learning About Statistical Methods

While writing the code, I realized I didn't fully understand what the methods were doing. My Jupyter Notebook code worked well because I either copied code from lectures and adjusted it, or the AI tools provided suggestions. I therefore dedicated time to become more familiar with the methods themselves—focusing not just on the Python implementation, but on what they actually analyze, when to use them, and how to interpret their results. The AI provided a summary of my work and suggested Python code, which I used to verify my own implementation.

For example:

Seasonally Adjusted Temperature Variations (SATV) uses a Discrete Cosine Transform (DCT) high-pass filter to remove slow-changing trends or noise, helping highlight short-term fluctuations. It then applies Statistical Process Control (SPC) limits to detect outliers.

Local Outlier Factor (LOF) is an unsupervised algorithm that detects anomalies by finding points with much lower density than their neighbors—useful for spotting isolated outliers.